

- **Creating the model**

- -> getting the data in the right form / shape is the hardest part
- -> this is creating the model
- **-> creating the model using keras**
 - -> keras is used to create neural networks in Python
 - -> each line of this code is a layer in the neural network
- **-> using keras to pass through the layers sequentially -> i.e one by one**
 - -> taking the 28x28 structure which stores the images and flattening them into a neural network
 - -> the second line, layer 2 -> is a dense layer
 - -> all of the neurones in the previous layer are connected to all of the neurones in the next
 - -> and there is an activation function being used -> to normalise those values and make the algorithm easier to run
 - -> sigmoid, tanh -> there are different activation layers we can use
 - **-> then the last layer**
 - -> there are 10 neurones in this layer
 - -> there are 10 different items of clothing which the images inputted into the model could be
 - -> softmax <- to make sure that the values for each of those nodes sum to one
 - -> each of the values at those nodes are the probabilities that that node will be that specific item of clothing

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)), # input layer (1)
    keras.layers.Dense(128, activation='relu'), # hidden layer (2)
    keras.layers.Dense(10, activation='softmax') # output layer (3)
])
```

- -> this is the architecture of the neural network
 - -> the number of neurones at each layer
 - -> the activation function
 - -> the type of connections
- **-> compiling the model**
 - -> picking the optimiser, the loss function and the metrics which we're going to use
 - the function we're going to use to optimise -> and the value we want to increase (the accuracy of the model)
 - -> the optimiser is the function which is going to perform the gradient descent
 - -> hyperparameters are things we can change -> e.g the number of nodes in the model or the activation function
 - -> hyperparameter tuning is changing these values to optimise the model
 - -> compiling the model
- **-> training the model**
 - -> fitting the model to the training data
 - -> this is the same as training the model
 - -> during each stage of the training, in this case it's printing out the success rates of each iteration as the model
 - -> it's being ran on the cloud -> google collaboratory
 - -> the amount of resources it takes to train the model
 - -> the model has an accuracy when it's being trained -> the actual accuracy of the model is how it behaves when it's being ran on the data which it wasn't trained on
 - -> (below) this was the code used to train the model

```
4423680/4422102 [=====] - 0s 0us/step

[3] model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)), # input layer (1)
    keras.layers.Dense(128, activation='relu'), # hidden layer (2)
    keras.layers.Dense(10, activation='softmax') # output layer (3)
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(train_images, train_labels, epochs=10)

... Train on 60000 samples
```

- -> the model has been fit
- -> **finding the true accuracy of the model**
- -> testing it on the testing data

```
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=1)

print('Test accuracy:', test_acc)
```

- -> verbose dictates how much data is being printed out when we run the model
- -> overfitting -> this is when the accuracy of the model is high on the training data and low on the testing data <- because it wasn't being given enough data when training the model
 - -> we need to make sure that the model generalises properly
 - -> overfitting and hyper-parameter tuning
 - -> evaluating the model is how we print its accuracy
 - -> the accuracy can be suspiciously high, in which case it may have been overfit
 - -> less epochs can improve the accuracy of the model -> more epochs means less information per epoch which means the model can be overfit and the accuracy can be too high
 - -> then when it comes to testing the model on data it hasn't seen before the accuracy drops
- -> **using the model to make predictions**

```
predictions = model.predict([test_images[0]])
```

- -> `test_images[0]` <- this is an array which you want the model to predict
- -> it is an array of different pixels for an image
- -> and we are running the model on it
- -> the predictions it returns is an array
 - -> that array is an array of probabilities
 - -> then using numpy to return the maximum value of that list -> aka the highest probability
 - -> which means it's a boot
- -> **what it's done is**
 - -> the image is represented by an array of pixels
 - -> the model was ran on it
 - -> then it's returned an array of probabilities

- -> each element in that array contains the probability that the image is a picture of a different item of clothing
- -> the largest number (max) in that output array is the item of clothing which that image is most likely to contain <- in other words the prediction for the model
- **-> making predictions on any entry that we want**
 - -> making predictions on the model
 - -> he then he's running the code in another cell where you input a random number
 - -> (below) this is the code which does this

```

COLOR = 'white'
plt.rcParams['text.color'] = COLOR
plt.rcParams['axes.labelcolor'] = COLOR

def predict(model, image, correct_label):
    class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
                   'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
    prediction = model.predict(np.array([image]))
    predicted_class = class_names[np.argmax(prediction)]

    show_image(image, class_names[correct_label], predicted_class)

def show_image(img, label, guess):
    plt.figure()
    plt.imshow(img, cmap=plt.cm.binary)
    plt.title("Expected: " + label)
    plt.xlabel("Guess: " + guess)
    plt.colorbar()
    plt.grid(False)
    plt.show()

def get_number():
    while True:
        num = input("Pick a number: ")
        if num.isdigit():
            num = int(num)
            if 0 <= num <= 1000:
                return int(num)
        else:
            print("Try again...")

num = get_number()
image = test_images[num]
label = test_labels[num]
predict(model, image, label)

```

• Next is

- -> convolutional neural networks
- -> deep computer vision / object recognition and detection
- -> building a model of dense layers

