- ***How convolutional networks / layers work***
  - -> in this example there is an image
  - -> the image is made of pixels which can be divided into a grid
  - ***properties of convolutional maps***
    - -> a feature map / filter
    - -> properties of the convolutional map
      - the imput size <- the number of filters and the same size of the fileters
      - -> each of the filters is looking for a pattern in the image (e.g the presence of paws)
    - -> the filters are found by the convolutional neural network
  - ***how we look for filters in the image***
    - -> one filter might look for a diagonal line, a horizontal line or a straight line
    - -> there are 32 different ones in this example
    - -> there is a filter size
    - -> 3x3 spots in the image
    - -> defining how closely those spots match
    - **-> doing the dot product of the filter with the image we have**
      - **-> the component of that feature in the direction of this image**
      - **-> and each filter is looking for the presence of a different feature in the image**
      - -> you can store the outputs of the filters (different dot products with the image we have) in a matrix
      - -> this is a response map
      - -> there are 3x3 samples -> r, g, b

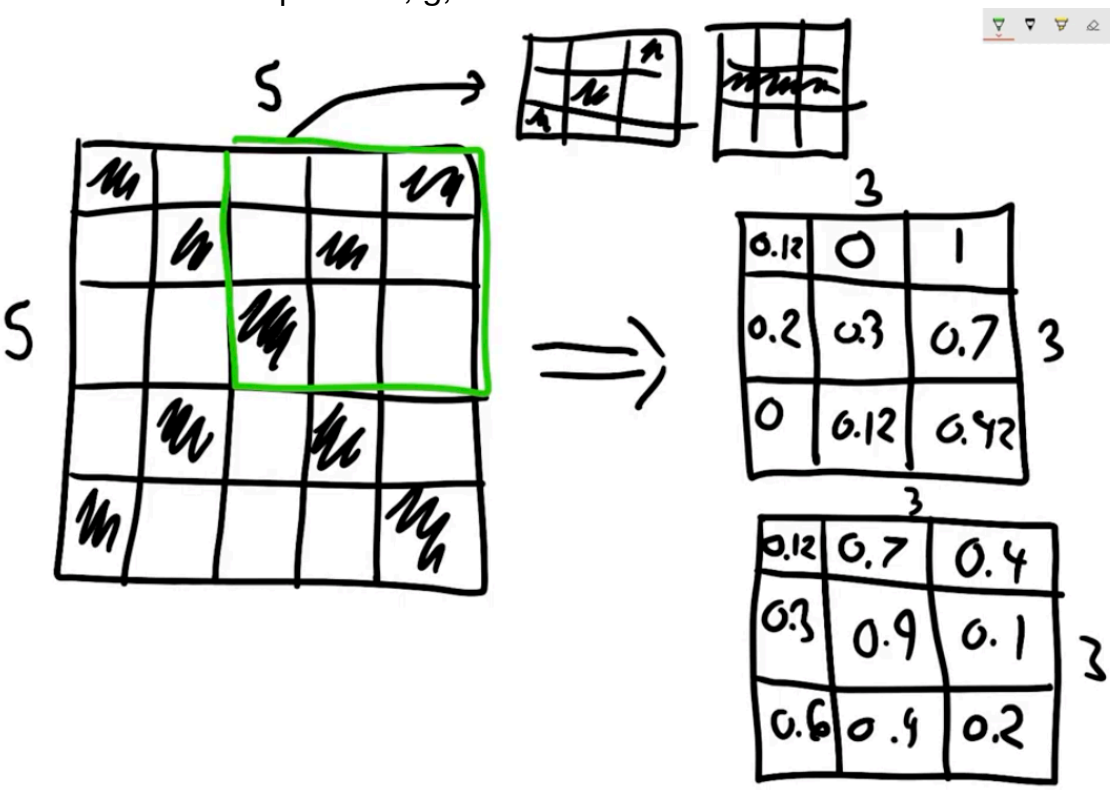

      - -> it's telling you about the presence of features in the image
      - -> this is done for many different features -> computationally heavy
- ***Pooling***
  - -> there is an output feature map
  - -> running the calculations is computationally heavy
  - -> looking at for example - the presence of different line types, to the combinations of different lines which we have -> pooling the information about the different lines which we have
    - -> we are building up to this, first from determining if lines are present - and then to finding correlations between those different lines which are present -> to identify the presence of e.g eyes in an image

- ***-> you can add padding around the image***
  - -> so that you can draw a square which makes the pixels on the edge in the centre
  - -> this allows you to include the pixels which are on the edge of the image

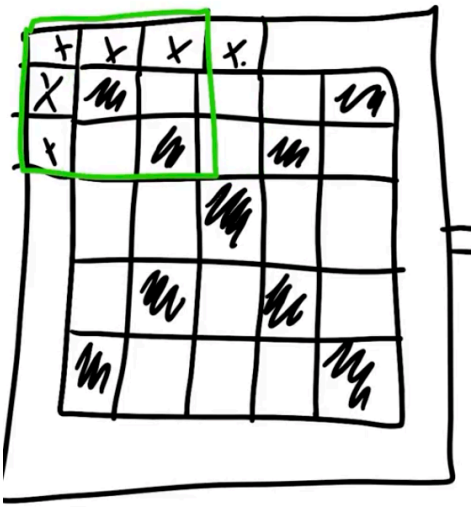

- ***-> stride***
  - -> this is how much the sample box is moved each time
  - -> for example 2px -> moving the sample box by 2px each time the calculations are ran
- ***-> a pooling operation***
  - -> min
  - -> max
  - -> average
  - -> taking specific values from a sample of the output feature
  - -> sampling 2x2 areas of the feature map
    - -> taking the min / max / average values of those points
  - ***-> pooling is taking an average from multiple filters***
    - -> the filter contains a matrix of dot products for a different feature with the pixels int he image
    - -> run that multiple times with different filters that check for the lines
    - -> then take the average of them and put them into another matrix -> this is pooling
    - -> this is done with a stride of 1 (in this example)
    - -> max pooling can be done
    - -> whether the feature exists
  - ***-> min / max / average tell you different things***
    - -> it's a way of combining the different matrices
    - -> the average presence of the feature, does the feature exist, does the feature not exist
    - -> the three main properties of each convolutional layer are the input size, the number of filters, and the sample size of the filters