

- **Creating a convolutional neural network**

- -> this is using keras
- -> also a dataset of images -> which contains matrices
- -> looking up what you need in the documentation
- -> there are different objects in the dataset and 60,000 images
- **Loading in the data**

```
tensorflow is already loaded. Please restart the runtime to change versions.

# LOAD AND SPLIT DATASET
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0


class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
```

- -> then normalising the data into training and test images
 - -> dividing them by 255 to make sure that they are normalised and the algorithms run properly
 - -> then the names of the classes
- -> then he is changing the indices of different values in the cell -> which returns different images in the dataset

```
# Let's look at a one image
IMG_INDEX = 7 # change this to look at other images

plt.imshow(train_images[IMG_INDEX], cmap=plt.cm.binary)
plt.xlabel(class_names[train_labels[IMG_INDEX][0]])
plt.show()
```



- **Making the convolutional neural network**

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

- -> cnn <- convolutional neural network
- -> stacking convolutional layers / max pooling / min pooling layers
- -> each of the lines of code are a layer in the network
 - -> max pooling layers
 - -> layers to reduce the dimensionality
 - -> you define the amount of filters, the elements per filter, the activation function ->

to normalise the dot products after running the filter, then the input shape -> the rows by columns of the input to the model

```
model.summary() # let's have a look at our model so far
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_4 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_5 (Conv2D)	(None, 4, 4, 64)	36928

Total params: 56,320
Trainable params: 56,320
Non-trainable params: 0

- -> it's 30x30x32 -> because there are two pixels of padding
- -> shrinking the shape by a factor of 2
- -> taking 64 filters
- -> max pooling
- -> this is the stack of convolution and max pooling layers

○ Adding the dense layers

- -> this is a convolution base which extracts the features from the image
- -> then we add the dense layers which take those different features and combine them
 - -> these combination of features can be used to classify the object
- -> to add the dense layers
 - -> `model.add(layers.Flatten())` <- moving them into one dimension
 - -> a 64 dense layer
- -> after adding these layers he has then printed out another summary of the model

```
model.add(layers.Flatten())  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(10))  
  
[ ] model.summary()
```

- -> 10 is the amount of classes which we are classifying it into
- -> the convolutional base / classifier

○ Training the model

- -> he's reducing the epochs to 4 -> it should be on 10 but this takes too long
- -> when the model is trained its accuracy is returned
- -> the loss function computes the cross entropy loss
- -> use `atom` / a categorical cross entropy loss function -> there are a lot of different loss

Note: This will take much longer than previous models!

```
model.compile(optimizer='adam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['accuracy'])  
  
history = model.fit(train_images, train_labels, epochs=4,  
                    validation_data=(test_images, test_labels))
```

functions and you can look them up depending on the context