


- there are three states
- in each state either the agent can stay or leave
- the agent needs to take the action
- **the agent wants to maximise the reward in the environment**
 - -> the number of states, actions and ways the agent can interact with the states
 - -> staying in the same state can give it a reward
 - -> the agent can interact with the environment which changes its state
- **the story of the agent**
 - -> the agent starts at state s1
 - -> it can either stay in the current state or move and receive a reward
 - -> he's making a matrix of all of the different possibilities and what the rewards are for each of them
 - -> looking at the table one time for each state
 - -> depending on where the agent starts



- -> it can be trapped in local minima -> when it could be receiving a much larger reward elsewhere
- -> exploring the environment and observing the rewards is how the Q-table is filled
 - you can end up with millions of these languages
 - once we have the table, then the model can produce the optimal actions depending on the state the agent is in
 - -> this is not based off of previous experiences -> where to go next is based off of the rewards of going there (which are independent of which state the agent came from)
 - -> taking random actions and being able to explore the environment more freely
- **-> learning the Q-table <- the table of rewards for different actions**
 - -> you need the agent to go after different actions which haven't been tried before in order to see what the rewards are
 - -> the model can randomly take an action, or pick one based off of what the greatest reward would be <- the actual algorithm does a mixture of both
 - -> if this is not done then it can get stuck in a local maxima - missing out the reward from another state
 - -> when it gets into a new state, it keeps updating the current environment

$$Q[state, action] = Q[state, action] + \alpha * (reward + \gamma * \max(Q[newState, :]) - Q[state, action])$$

- -> α is the learning rate
- -> γ is for the discount factor
 - -> the balance between taking a random action and one which would reap the highest rewards based off of where the agent was
 - -> the move action had to have a higher reward value than having it stay
 - this factor is how much the Q-values are allowed to be updated by after every single action
 - -> the learning rate means it will update slower
 - -> the value which is added is positive or negative
 - -> it's asking what the maximum reward given the new state which it's just moved into is
- -> factoring in the reward to determine the best place to move into
- -> the transition states
- -> subtracting the states and actions to prevent the agent from getting stuck in the same state
- -> the learning rate tells you how much you can update each learning value by
 - -> the balance between the future actions and the current one